

# MIS 583 Assignment 5: YOLO Object Detection on PASCAL VOC

Before we start, please put your name and SID in following format:

: LASTNAME Firstname, ?00000000 // e.g.) 李晨愷 M114020035

## Your Answer:

Hi I'm XXX, XXXXXXXXXXXX.

## Google Colab Setup

Next we need to run a few commands to set up our environment on Google Colab. If you are running this notebook on a local machine you can skip this section.

Run the following cell to mount your Google Drive. Follow the link, sign in to your Google account (the same account you used to store this notebook!) and copy the authorization code into the text box that appears below.

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

## How to Get Data

請先到共用雲端硬碟將檔案 V0Cdevkit\_2007.zip，建立捷徑到自己的雲端硬碟中。

### 操作步驟

1. 點開雲端連結
2. 點選右上角「新增雲端硬碟捷徑」
3. 點選「我的雲端硬碟」
4. 點選「新增捷徑」

完成以上流程會在你的雲端硬碟中建立一個檔案的捷徑，接著我們在colab中取得權限即可使用。

## Unzip Data

解壓縮 V0Cdevkit\_2007.zip

- V0C2007 : 包含了train/val的所有圖片
- V0C2007test : 包含了test的所有圖片

其中 train 的圖片 3756 張， val 的圖片 1255 張， test 的圖片 4950 張。

注意: 若有另外設定存放在雲端硬碟中的路徑，請記得本處路徑也須做更動。

**Notice: Please put "VOCdevkit\_2007" folder under data folder.**

```
In [ ]: !unzip -qq ./drive/MyDrive/VOCdevkit_2007.zip
```

## Import package

```
In [1]: import os
import random

import cv2
import numpy as np

import csv

import torch
from torch.utils.data import DataLoader
from torchvision import models

from src.resnet_yolo import resnet50
from yolo_loss import YoloLoss
from src.dataset import VocDetectorDataset
from src.eval_voc import evaluate, test_evaluate
from src.predict import predict_image
from src.config import VOC_CLASSES, COLORS
from kaggle_submission import write_csv

import matplotlib.pyplot as plt
import collections

%matplotlib inline
%load_ext autoreload
%autoreload 2
```

## Initialization

```
In [2]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [3]: # YOLO network hyperparameters
B = 2 # number of bounding box predictions per cell
S = 14 # width/height of network output grid (larger than 7x7 from paper)
```

To implement Yolo we will rely on a pretrained classifier as the backbone for our detection network. PyTorch offers a variety of models which are pretrained on ImageNet in the `torchvision.models` package. In particular, we will use the ResNet50 architecture as a base for our detector. This is different from the base architecture in the Yolo paper and also results in a different output grid size (14x14 instead of 7x7).

Models are typically pretrained on ImageNet since the dataset is very large (> 1 million images) and widely used. The pretrained model provides a very useful weight initialization for our detector, so that the network is able to learn quickly and effectively.

```
In [4]: load_network_path = None #'checkpoints/best_detector.pth'
pretrained = True

# use to load a previously trained network
if load_network_path is not None:
    print('Loading saved network from {}'.format(load_network_path))
    net = resnet50().to(device)
    net.load_state_dict(torch.load(load_network_path))
else:
    print('Load pre-trained model')
    net = resnet50(pretrained=pretrained).to(device)
```

Load pre-trained model

```
/home/vllab/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
```

```
warnings.warn(
/home/vllab/anaconda3/lib/python3.11/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet50_Weights.IMAGE NET1K_V1`. You can also use `weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights.
```

```
warnings.warn(msg)
```

```
In [5]: learning_rate = 0.001
num_epochs = 50 #50
batch_size = 10 #24

# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5
```

## Reading Pascal Data

The train dataset loader also using a variety of data augmentation techniques including random shift, scaling, crop, and flips. Data augmentation is slightly more complicated for detection datasets since the bounding box annotations must be kept consistent throughout the transformations.

Since the output of the detector network we train is an  $S \times S \times (B \times 5 + C)$ , we use an encoder to convert the original bounding box coordinates into relative grid bounding box coordinates corresponding to the expected output. We also use a decoder which allows us to convert the opposite direction into image coordinate bounding boxes.

**Notice: Please put "VOCdevkit\_2007" folder under data folder.**

```
In [6]: file_root_train = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
annotation_file_train = 'data/voc2007train.txt'

train_dataset = VocDetectorDataset(root_img_dir=file_root_train, dataset_f
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
print('Loaded %d train images' % len(train_dataset))
```

Initializing dataset

Loaded 3756 train images

```
In [7]: file_root_val = 'data/VOCdevkit_2007/VOC2007/JPEGImages/'
        annotation_file_val = 'data/voc2007val.txt'

        val_dataset = VocDetectorDataset(root_img_dir=file_root_val, dataset_file=
        val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, n
        print('Loaded %d val images' % len(val_dataset))

        Initializing dataset
        Loaded 1255 val images
```

```
In [8]: data = train_dataset[0]
```

## Set up training tools

```
In [9]: criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
        optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=
```

## Train detector

```

In [10]: best_val_loss = np.inf
learning_rate = 1e-3
for epoch in range(num_epochs):
    torch.cuda.empty_cache()
    net.train()

    # Update learning rate late in training
    if epoch == 30 or epoch == 40:
        learning_rate /= 10.0

    for param_group in optimizer.param_groups:
        param_group['lr'] = learning_rate

    print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
    print('Learning Rate for this epoch: {}'.format(learning_rate))

    total_loss = collections.defaultdict(int)

    for i, data in enumerate(train_loader):
        data = (item.to(device) for item in data)
        images, target_boxes, target_cls, has_object_map = data
        pred = net(images)
        loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
        for key in loss_dict:
            total_loss[key] += loss_dict[key].item()

        optimizer.zero_grad()
        loss_dict['total_loss'].backward()
        optimizer.step()

        if (i+1) % 50 == 0:
            outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1), num_epochs, i+1, 50)
            outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for key, val in loss_dict.items())
            print(outstring)

    # evaluate the network on the val data
    if (epoch + 1) % 5 == 0:
        val_aps = evaluate(net, val_dataset_file=annotation_file_val, img_loader=val_loader)
        print(epoch, val_aps)
    with torch.no_grad():
        val_loss = 0.0
        net.eval()
        for i, data in enumerate(val_loader):
            data = (item.to(device) for item in data)
            images, target_boxes, target_cls, has_object_map = data

            pred = net(images)
            loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
            val_loss += loss_dict['total_loss'].item()
        val_loss /= len(val_loader)

    if best_val_loss > val_loss:
        best_val_loss = val_loss
        print('Updating best val loss: %.5f' % best_val_loss)
        torch.save(net.state_dict(), 'checkpoints/best_detector.pth')

    if (epoch+1) in [5, 10, 20, 30, 40]:
        torch.save(net.state_dict(), 'checkpoints/detector_epoch_%d.pth' % (epoch+1))

    torch.save(net.state_dict(), 'checkpoints/detector.pth')

```

Starting epoch 1 / 50

Learning Rate for this epoch: 0.001

Epoch [1/50], Iter [50/376], Loss: total=25.513, reg=4.523, containing\_obj=0.325, no\_obj=12.562, cls=8.102

Epoch [1/50], Iter [100/376], Loss: total=17.515, reg=4.121, containing\_obj=0.410, no\_obj=6.510, cls=6.474

Epoch [1/50], Iter [150/376], Loss: total=14.232, reg=3.711, containing\_obj=0.471, no\_obj=4.446, cls=5.605

Epoch [1/50], Iter [200/376], Loss: total=12.407, reg=3.465, containing\_obj=0.498, no\_obj=3.397, cls=5.048

Epoch [1/50], Iter [250/376], Loss: total=11.233, reg=3.296, containing\_obj=0.532, no\_obj=2.759, cls=4.646

Epoch [1/50], Iter [300/376], Loss: total=10.404, reg=3.166, containing\_obj=0.558, no\_obj=2.329, cls=4.351

Epoch [1/50], Iter [350/376], Loss: total=9.768, reg=3.063, containing\_obj=0.581, no\_obj=2.019, cls=4.105

Updating best val loss: 5.75227

Starting epoch 2 / 50

Learning Rate for this epoch: 0.001

Epoch [2/50], Iter [50/376], Loss: total=5.409, reg=2.228, containing\_obj=0.739, no\_obj=0.134, cls=2.308

Epoch [2/50], Iter [100/376], Loss: total=5.431, reg=2.226, containing\_obj=0.751, no\_obj=0.128, cls=2.325

Epoch [2/50], Iter [150/376], Loss: total=5.440, reg=2.225, containing\_obj=0.762, no\_obj=0.123, cls=2.329

Epoch [2/50], Iter [200/376], Loss: total=5.354, reg=2.207, containing\_obj=0.770, no\_obj=0.119, cls=2.258

Epoch [2/50], Iter [250/376], Loss: total=5.301, reg=2.206, containing\_obj=0.776, no\_obj=0.115, cls=2.204

Epoch [2/50], Iter [300/376], Loss: total=5.243, reg=2.196, containing\_obj=0.772, no\_obj=0.111, cls=2.164

Epoch [2/50], Iter [350/376], Loss: total=5.240, reg=2.203, containing\_obj=0.777, no\_obj=0.108, cls=2.153

Updating best val loss: 5.20640

Starting epoch 3 / 50

Learning Rate for this epoch: 0.001

Epoch [3/50], Iter [50/376], Loss: total=4.688, reg=2.073, containing\_obj=0.792, no\_obj=0.082, cls=1.741

Epoch [3/50], Iter [100/376], Loss: total=4.797, reg=2.102, containing\_obj=0.807, no\_obj=0.080, cls=1.808

Epoch [3/50], Iter [150/376], Loss: total=4.774, reg=2.064, containing\_obj=0.811, no\_obj=0.078, cls=1.820

Epoch [3/50], Iter [200/376], Loss: total=4.808, reg=2.090, containing\_obj=0.816, no\_obj=0.077, cls=1.825

Epoch [3/50], Iter [250/376], Loss: total=4.780, reg=2.077, containing\_obj=0.822, no\_obj=0.076, cls=1.805

Epoch [3/50], Iter [300/376], Loss: total=4.778, reg=2.085, containing\_obj=0.829, no\_obj=0.074, cls=1.790

Epoch [3/50], Iter [350/376], Loss: total=4.768, reg=2.080, containing\_obj=0.830, no\_obj=0.073, cls=1.784

Updating best val loss: 4.87400

Starting epoch 4 / 50

Learning Rate for this epoch: 0.001

Epoch [4/50], Iter [50/376], Loss: total=4.495, reg=1.937, containing\_obj=0.821, no\_obj=0.065, cls=1.673

Epoch [4/50], Iter [100/376], Loss: total=4.437, reg=1.950, containing\_obj

```
j=0.819, no_obj=0.064, cls=1.603
Epoch [4/50], Iter [150/376], Loss: total=4.418, reg=1.944, containing_obj
j=0.830, no_obj=0.064, cls=1.580
Epoch [4/50], Iter [200/376], Loss: total=4.497, reg=1.991, containing_obj
j=0.842, no_obj=0.064, cls=1.601
Epoch [4/50], Iter [250/376], Loss: total=4.436, reg=1.964, containing_obj
j=0.842, no_obj=0.064, cls=1.566
Epoch [4/50], Iter [300/376], Loss: total=4.426, reg=1.960, containing_obj
j=0.836, no_obj=0.064, cls=1.567
Epoch [4/50], Iter [350/376], Loss: total=4.435, reg=1.969, containing_obj
j=0.846, no_obj=0.064, cls=1.557
Updating best val loss: 4.67710
```

```
Starting epoch 5 / 50
Learning Rate for this epoch: 0.001
Epoch [5/50], Iter [50/376], Loss: total=4.580, reg=2.039, containing_obj
=0.905, no_obj=0.066, cls=1.570
Epoch [5/50], Iter [100/376], Loss: total=4.488, reg=1.972, containing_obj
j=0.884, no_obj=0.067, cls=1.564
Epoch [5/50], Iter [150/376], Loss: total=4.411, reg=1.940, containing_obj
j=0.877, no_obj=0.068, cls=1.526
Epoch [5/50], Iter [200/376], Loss: total=4.375, reg=1.932, containing_obj
j=0.869, no_obj=0.069, cls=1.505
Epoch [5/50], Iter [250/376], Loss: total=4.342, reg=1.923, containing_obj
j=0.864, no_obj=0.070, cls=1.485
Epoch [5/50], Iter [300/376], Loss: total=4.328, reg=1.917, containing_obj
j=0.867, no_obj=0.071, cls=1.473
Epoch [5/50], Iter [350/376], Loss: total=4.300, reg=1.903, containing_obj
j=0.860, no_obj=0.072, cls=1.465
---Evaluate model on test samples---
100%|██████████| 1255/1255 [00:27<00:00, 46.27it/s]
```

```
---class aeroplane ap 0.025---
---class bicycle ap 0.0---
---class bird ap 0.0--- (no predictions for this class)
---class boat ap 0.0--- (no predictions for this class)
---class bottle ap 0.0--- (no predictions for this class)
---class bus ap 0.0--- (no predictions for this class)
---class car ap 0.1964678090677125---
---class cat ap 0.0--- (no predictions for this class)
---class chair ap 0.005---
---class cow ap 0.0--- (no predictions for this class)
---class diningtable ap 0.0--- (no predictions for this class)
---class dog ap 0.0--- (no predictions for this class)
---class horse ap 0.0681063122923588---
---class motorbike ap 0.0--- (no predictions for this class)
---class person ap 0.0835911979597899---
---class pottedplant ap 0.0--- (no predictions for this class)
---class sheep ap 0.0---
---class sofa ap 0.0--- (no predictions for this class)
---class train ap 0.0--- (no predictions for this class)
---class tvmonitor ap 0.0---
---map 0.018908265965993064---
4 [0.025, 0.0, 0.0, 0.0, 0.0, 0.0, 0.1964678090677125, 0.0, 0.005, 0.0,
0.0, 0.0, 0.0681063122923588, 0.0, 0.0835911979597899, 0.0, 0.0, 0.0, 0.
0, 0.0]
Updating best val loss: 4.48595
```

```
Starting epoch 6 / 50
Learning Rate for this epoch: 0.001
Epoch [6/50], Iter [50/376], Loss: total=3.969, reg=1.729, containing_obj
=0.831, no_obj=0.083, cls=1.325
Epoch [6/50], Iter [100/376], Loss: total=4.160, reg=1.836, containing_ob
j=0.851, no_obj=0.085, cls=1.388
Epoch [6/50], Iter [150/376], Loss: total=4.129, reg=1.828, containing_ob
j=0.846, no_obj=0.086, cls=1.369
Epoch [6/50], Iter [200/376], Loss: total=4.145, reg=1.849, containing_ob
j=0.848, no_obj=0.088, cls=1.360
Epoch [6/50], Iter [250/376], Loss: total=4.131, reg=1.830, containing_ob
j=0.846, no_obj=0.089, cls=1.366
Epoch [6/50], Iter [300/376], Loss: total=4.110, reg=1.817, containing_ob
j=0.844, no_obj=0.090, cls=1.359
Epoch [6/50], Iter [350/376], Loss: total=4.098, reg=1.818, containing_ob
j=0.841, no_obj=0.090, cls=1.348
Updating best val loss: 4.29943
```

```
Starting epoch 7 / 50
Learning Rate for this epoch: 0.001
Epoch [7/50], Iter [50/376], Loss: total=4.010, reg=1.779, containing_obj
=0.841, no_obj=0.092, cls=1.298
Epoch [7/50], Iter [100/376], Loss: total=3.916, reg=1.727, containing_ob
j=0.825, no_obj=0.093, cls=1.271
Epoch [7/50], Iter [150/376], Loss: total=3.935, reg=1.743, containing_ob
j=0.821, no_obj=0.094, cls=1.277
Epoch [7/50], Iter [200/376], Loss: total=3.884, reg=1.726, containing_ob
j=0.809, no_obj=0.096, cls=1.254
Epoch [7/50], Iter [250/376], Loss: total=3.947, reg=1.762, containing_ob
j=0.831, no_obj=0.096, cls=1.258
Epoch [7/50], Iter [300/376], Loss: total=3.910, reg=1.742, containing_ob
j=0.826, no_obj=0.096, cls=1.246
Epoch [7/50], Iter [350/376], Loss: total=3.927, reg=1.755, containing_ob
j=0.832, no_obj=0.097, cls=1.242
Updating best val loss: 4.18828
```



Starting epoch 8 / 50  
Learning Rate for this epoch: 0.001  
Epoch [8/50], Iter [50/376], Loss: total=3.873, reg=1.712, containing\_obj=0.868, no\_obj=0.098, cls=1.195  
Epoch [8/50], Iter [100/376], Loss: total=3.920, reg=1.756, containing\_obj=0.869, no\_obj=0.099, cls=1.195  
Epoch [8/50], Iter [150/376], Loss: total=3.947, reg=1.782, containing\_obj=0.877, no\_obj=0.099, cls=1.189  
Epoch [8/50], Iter [200/376], Loss: total=3.907, reg=1.755, containing\_obj=0.862, no\_obj=0.102, cls=1.189  
Epoch [8/50], Iter [250/376], Loss: total=3.865, reg=1.731, containing\_obj=0.852, no\_obj=0.102, cls=1.180  
Epoch [8/50], Iter [300/376], Loss: total=3.812, reg=1.711, containing\_obj=0.837, no\_obj=0.103, cls=1.161  
Epoch [8/50], Iter [350/376], Loss: total=3.778, reg=1.697, containing\_obj=0.828, no\_obj=0.103, cls=1.149  
Updating best val loss: 3.97744

Starting epoch 9 / 50  
Learning Rate for this epoch: 0.001  
Epoch [9/50], Iter [50/376], Loss: total=3.627, reg=1.623, containing\_obj=0.848, no\_obj=0.112, cls=1.043  
Epoch [9/50], Iter [100/376], Loss: total=3.679, reg=1.634, containing\_obj=0.831, no\_obj=0.111, cls=1.103  
Epoch [9/50], Iter [150/376], Loss: total=3.670, reg=1.641, containing\_obj=0.831, no\_obj=0.108, cls=1.090  
Epoch [9/50], Iter [200/376], Loss: total=3.622, reg=1.628, containing\_obj=0.817, no\_obj=0.110, cls=1.066  
Epoch [9/50], Iter [250/376], Loss: total=3.552, reg=1.592, containing\_obj=0.803, no\_obj=0.111, cls=1.046  
Epoch [9/50], Iter [300/376], Loss: total=3.544, reg=1.591, containing\_obj=0.801, no\_obj=0.112, cls=1.041  
Epoch [9/50], Iter [350/376], Loss: total=3.568, reg=1.609, containing\_obj=0.807, no\_obj=0.111, cls=1.041  
Updating best val loss: 3.89062

Starting epoch 10 / 50  
Learning Rate for this epoch: 0.001  
Epoch [10/50], Iter [50/376], Loss: total=3.527, reg=1.592, containing\_obj=0.808, no\_obj=0.112, cls=1.015  
Epoch [10/50], Iter [100/376], Loss: total=3.509, reg=1.592, containing\_obj=0.826, no\_obj=0.113, cls=0.978  
Epoch [10/50], Iter [150/376], Loss: total=3.489, reg=1.586, containing\_obj=0.813, no\_obj=0.113, cls=0.977  
Epoch [10/50], Iter [200/376], Loss: total=3.509, reg=1.590, containing\_obj=0.815, no\_obj=0.111, cls=0.993  
Epoch [10/50], Iter [250/376], Loss: total=3.465, reg=1.571, containing\_obj=0.812, no\_obj=0.111, cls=0.971  
Epoch [10/50], Iter [300/376], Loss: total=3.482, reg=1.579, containing\_obj=0.811, no\_obj=0.112, cls=0.980  
Epoch [10/50], Iter [350/376], Loss: total=3.475, reg=1.579, containing\_obj=0.810, no\_obj=0.113, cls=0.974  
---Evaluate model on test samples---  
100%|██████████| 1255/1255 [00:31<00:00, 39.58it/s]

```
---class aeroplane ap 0.3643968402920609---
---class bicycle ap 0.1922373583939121---
---class bird ap 0.1549462249066992---
---class boat ap 0.06230938028927922---
---class bottle ap 0.0--- (no predictions for this class)
---class bus ap 0.045454545454545456---
---class car ap 0.20934365069154598---
---class cat ap 0.16304916193681487---
---class chair ap 0.14782833580839763---
---class cow ap 0.002898550724637681---
---class diningtable ap 0.015873015873015872---
---class dog ap 0.1223697140146381---
---class horse ap 0.3135874682544261---
---class motorbike ap 0.1006728778467909---
---class person ap 0.22540235175931725---
---class pottedplant ap 0.018017344497607654---
---class sheep ap 0.04662058371735791---
---class sofa ap 0.0--- (no predictions for this class)
---class train ap 0.3447655374268278---
---class tvmonitor ap 0.21836505100352627---
---map 0.13740689964457004---
9 [0.3643968402920609, 0.1922373583939121, 0.1549462249066992, 0.06230938
028927922, 0.0, 0.045454545454545456, 0.20934365069154598, 0.163049161936
81487, 0.14782833580839763, 0.002898550724637681, 0.015873015873015872,
0.1223697140146381, 0.3135874682544261, 0.1006728778467909, 0.22540235175
931725, 0.018017344497607654, 0.04662058371735791, 0.0, 0.344765537426827
8, 0.21836505100352627]
Updating best val loss: 3.73310
```

Starting epoch 11 / 50

Learning Rate for this epoch: 0.001

Epoch [11/50], Iter [50/376], Loss: total=3.599, reg=1.707, containing\_obj=0.805, no\_obj=0.121, cls=0.966

Epoch [11/50], Iter [100/376], Loss: total=3.495, reg=1.637, containing\_obj=0.788, no\_obj=0.127, cls=0.943

Epoch [11/50], Iter [150/376], Loss: total=3.510, reg=1.636, containing\_obj=0.798, no\_obj=0.127, cls=0.950

Epoch [11/50], Iter [200/376], Loss: total=3.485, reg=1.634, containing\_obj=0.799, no\_obj=0.125, cls=0.927

Epoch [11/50], Iter [250/376], Loss: total=3.441, reg=1.607, containing\_obj=0.798, no\_obj=0.125, cls=0.910

Epoch [11/50], Iter [300/376], Loss: total=3.406, reg=1.582, containing\_obj=0.796, no\_obj=0.125, cls=0.902

Epoch [11/50], Iter [350/376], Loss: total=3.377, reg=1.570, containing\_obj=0.788, no\_obj=0.127, cls=0.892

Updating best val loss: 3.68438

Starting epoch 12 / 50

Learning Rate for this epoch: 0.001

Epoch [12/50], Iter [50/376], Loss: total=3.250, reg=1.551, containing\_obj=0.788, no\_obj=0.122, cls=0.788

Epoch [12/50], Iter [100/376], Loss: total=3.211, reg=1.493, containing\_obj=0.790, no\_obj=0.128, cls=0.800

Epoch [12/50], Iter [150/376], Loss: total=3.230, reg=1.505, containing\_obj=0.791, no\_obj=0.129, cls=0.805

Epoch [12/50], Iter [200/376], Loss: total=3.224, reg=1.489, containing\_obj=0.797, no\_obj=0.130, cls=0.809

Epoch [12/50], Iter [250/376], Loss: total=3.246, reg=1.492, containing\_obj=0.805, no\_obj=0.131, cls=0.818

Epoch [12/50], Iter [300/376], Loss: total=3.238, reg=1.492, containing\_obj=0.796, no\_obj=0.133, cls=0.817

Epoch [12/50], Iter [350/376], Loss: total=3.229, reg=1.494, containing\_obj=0.794, no\_obj=0.132, cls=0.810  
Updating best val loss: 3.62757

Starting epoch 13 / 50

Learning Rate for this epoch: 0.001

Epoch [13/50], Iter [50/376], Loss: total=3.182, reg=1.532, containing\_obj=0.808, no\_obj=0.125, cls=0.717  
Epoch [13/50], Iter [100/376], Loss: total=3.209, reg=1.543, containing\_obj=0.794, no\_obj=0.130, cls=0.742  
Epoch [13/50], Iter [150/376], Loss: total=3.207, reg=1.536, containing\_obj=0.790, no\_obj=0.133, cls=0.747  
Epoch [13/50], Iter [200/376], Loss: total=3.182, reg=1.527, containing\_obj=0.790, no\_obj=0.134, cls=0.731  
Epoch [13/50], Iter [250/376], Loss: total=3.239, reg=1.541, containing\_obj=0.800, no\_obj=0.136, cls=0.762  
Epoch [13/50], Iter [300/376], Loss: total=3.227, reg=1.534, containing\_obj=0.794, no\_obj=0.137, cls=0.762  
Epoch [13/50], Iter [350/376], Loss: total=3.206, reg=1.525, containing\_obj=0.791, no\_obj=0.138, cls=0.752  
Updating best val loss: 3.44784

Starting epoch 14 / 50

Learning Rate for this epoch: 0.001

Epoch [14/50], Iter [50/376], Loss: total=3.159, reg=1.526, containing\_obj=0.787, no\_obj=0.142, cls=0.703  
Epoch [14/50], Iter [100/376], Loss: total=3.127, reg=1.486, containing\_obj=0.793, no\_obj=0.143, cls=0.704  
Epoch [14/50], Iter [150/376], Loss: total=3.075, reg=1.446, containing\_obj=0.777, no\_obj=0.147, cls=0.705  
Epoch [14/50], Iter [200/376], Loss: total=3.049, reg=1.437, containing\_obj=0.778, no\_obj=0.145, cls=0.689  
Epoch [14/50], Iter [250/376], Loss: total=3.059, reg=1.447, containing\_obj=0.773, no\_obj=0.145, cls=0.694  
Epoch [14/50], Iter [300/376], Loss: total=3.033, reg=1.428, containing\_obj=0.772, no\_obj=0.145, cls=0.688  
Epoch [14/50], Iter [350/376], Loss: total=3.029, reg=1.430, containing\_obj=0.771, no\_obj=0.145, cls=0.683

Starting epoch 15 / 50

Learning Rate for this epoch: 0.001

Epoch [15/50], Iter [50/376], Loss: total=3.091, reg=1.536, containing\_obj=0.751, no\_obj=0.142, cls=0.663  
Epoch [15/50], Iter [100/376], Loss: total=3.027, reg=1.480, containing\_obj=0.759, no\_obj=0.142, cls=0.646  
Epoch [15/50], Iter [150/376], Loss: total=2.975, reg=1.436, containing\_obj=0.752, no\_obj=0.146, cls=0.641  
Epoch [15/50], Iter [200/376], Loss: total=2.972, reg=1.431, containing\_obj=0.752, no\_obj=0.146, cls=0.642  
Epoch [15/50], Iter [250/376], Loss: total=2.990, reg=1.441, containing\_obj=0.759, no\_obj=0.147, cls=0.643  
Epoch [15/50], Iter [300/376], Loss: total=3.006, reg=1.449, containing\_obj=0.762, no\_obj=0.147, cls=0.648  
Epoch [15/50], Iter [350/376], Loss: total=2.976, reg=1.431, containing\_obj=0.756, no\_obj=0.147, cls=0.642

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:29<00:00, 42.57it/s]

```
---class aeroplane ap 0.4045504648081668---
---class bicycle ap 0.4463032691884042---
---class bird ap 0.2830405492191656---
---class boat ap 0.11454245410421258---
---class bottle ap 0.04894188859979564---
---class bus ap 0.30500097125097136---
---class car ap 0.4243559011497385---
---class cat ap 0.44621142185097856---
---class chair ap 0.10666651101138883---
---class cow ap 0.16640451957801572---
---class diningtable ap 0.0985134794658604---
---class dog ap 0.31808471724911386---
---class horse ap 0.3777354208684758---
---class motorbike ap 0.358902738910503---
---class person ap 0.3067608130342685---
---class pottedplant ap 0.04072315621518484---
---class sheep ap 0.07188683835163946---
---class sofa ap 0.20867516023574773---
---class train ap 0.5509435840315452---
---class tvmonitor ap 0.25565016639035504---
---map 0.26669470127567657---
14 [0.4045504648081668, 0.4463032691884042, 0.2830405492191656, 0.1145424
5410421258, 0.04894188859979564, 0.30500097125097136, 0.4243559011497385,
0.44621142185097856, 0.10666651101138883, 0.16640451957801572, 0.09851347
94658604, 0.31808471724911386, 0.3777354208684758, 0.358902738910503, 0.3
067608130342685, 0.04072315621518484, 0.07188683835163946, 0.208675160235
74773, 0.5509435840315452, 0.25565016639035504]
Updating best val loss: 3.35581
```

Starting epoch 16 / 50

Learning Rate for this epoch: 0.001

Epoch [16/50], Iter [50/376], Loss: total=2.722, reg=1.274, containing\_obj=0.713, no\_obj=0.154, cls=0.581

Epoch [16/50], Iter [100/376], Loss: total=2.781, reg=1.311, containing\_obj=0.722, no\_obj=0.151, cls=0.597

Epoch [16/50], Iter [150/376], Loss: total=2.795, reg=1.344, containing\_obj=0.720, no\_obj=0.151, cls=0.581

Epoch [16/50], Iter [200/376], Loss: total=2.807, reg=1.346, containing\_obj=0.730, no\_obj=0.151, cls=0.582

Epoch [16/50], Iter [250/376], Loss: total=2.822, reg=1.354, containing\_obj=0.731, no\_obj=0.152, cls=0.585

Epoch [16/50], Iter [300/376], Loss: total=2.834, reg=1.362, containing\_obj=0.736, no\_obj=0.152, cls=0.584

Epoch [16/50], Iter [350/376], Loss: total=2.841, reg=1.362, containing\_obj=0.737, no\_obj=0.153, cls=0.590

Updating best val loss: 3.28643

Starting epoch 17 / 50

Learning Rate for this epoch: 0.001

Epoch [17/50], Iter [50/376], Loss: total=2.765, reg=1.267, containing\_obj=0.760, no\_obj=0.163, cls=0.575

Epoch [17/50], Iter [100/376], Loss: total=2.696, reg=1.266, containing\_obj=0.724, no\_obj=0.161, cls=0.544

Epoch [17/50], Iter [150/376], Loss: total=2.768, reg=1.306, containing\_obj=0.752, no\_obj=0.160, cls=0.550

Epoch [17/50], Iter [200/376], Loss: total=2.757, reg=1.302, containing\_obj=0.744, no\_obj=0.158, cls=0.553

Epoch [17/50], Iter [250/376], Loss: total=2.748, reg=1.303, containing\_obj=0.740, no\_obj=0.158, cls=0.548

Epoch [17/50], Iter [300/376], Loss: total=2.770, reg=1.315, containing\_obj=0.743, no\_obj=0.157, cls=0.555

Epoch [17/50], Iter [350/376], Loss: total=2.783, reg=1.328, containing\_obj=0.746, no\_obj=0.157, cls=0.552

Starting epoch 18 / 50

Learning Rate for this epoch: 0.001

Epoch [18/50], Iter [50/376], Loss: total=2.674, reg=1.293, containing\_obj=0.717, no\_obj=0.156, cls=0.508

Epoch [18/50], Iter [100/376], Loss: total=2.732, reg=1.320, containing\_obj=0.741, no\_obj=0.157, cls=0.514

Epoch [18/50], Iter [150/376], Loss: total=2.682, reg=1.288, containing\_obj=0.735, no\_obj=0.158, cls=0.500

Epoch [18/50], Iter [200/376], Loss: total=2.645, reg=1.274, containing\_obj=0.721, no\_obj=0.160, cls=0.490

Epoch [18/50], Iter [250/376], Loss: total=2.670, reg=1.287, containing\_obj=0.724, no\_obj=0.162, cls=0.498

Epoch [18/50], Iter [300/376], Loss: total=2.716, reg=1.312, containing\_obj=0.738, no\_obj=0.161, cls=0.505

Epoch [18/50], Iter [350/376], Loss: total=2.714, reg=1.302, containing\_obj=0.739, no\_obj=0.161, cls=0.511

Updating best val loss: 3.23497

Starting epoch 19 / 50

Learning Rate for this epoch: 0.001

Epoch [19/50], Iter [50/376], Loss: total=2.824, reg=1.351, containing\_obj=0.792, no\_obj=0.154, cls=0.526

Epoch [19/50], Iter [100/376], Loss: total=2.783, reg=1.352, containing\_obj=0.757, no\_obj=0.159, cls=0.515

Epoch [19/50], Iter [150/376], Loss: total=2.750, reg=1.344, containing\_obj=0.749, no\_obj=0.163, cls=0.495

Epoch [19/50], Iter [200/376], Loss: total=2.736, reg=1.324, containing\_obj=0.745, no\_obj=0.165, cls=0.502

Epoch [19/50], Iter [250/376], Loss: total=2.711, reg=1.317, containing\_obj=0.728, no\_obj=0.166, cls=0.500

Epoch [19/50], Iter [300/376], Loss: total=2.716, reg=1.318, containing\_obj=0.731, no\_obj=0.166, cls=0.501

Epoch [19/50], Iter [350/376], Loss: total=2.691, reg=1.305, containing\_obj=0.723, no\_obj=0.166, cls=0.496

Updating best val loss: 3.18714

Starting epoch 20 / 50

Learning Rate for this epoch: 0.001

Epoch [20/50], Iter [50/376], Loss: total=2.646, reg=1.266, containing\_obj=0.739, no\_obj=0.164, cls=0.477

Epoch [20/50], Iter [100/376], Loss: total=2.598, reg=1.248, containing\_obj=0.733, no\_obj=0.168, cls=0.448

Epoch [20/50], Iter [150/376], Loss: total=2.627, reg=1.249, containing\_obj=0.738, no\_obj=0.171, cls=0.468

Epoch [20/50], Iter [200/376], Loss: total=2.622, reg=1.246, containing\_obj=0.733, no\_obj=0.170, cls=0.474

Epoch [20/50], Iter [250/376], Loss: total=2.624, reg=1.245, containing\_obj=0.726, no\_obj=0.168, cls=0.484

Epoch [20/50], Iter [300/376], Loss: total=2.606, reg=1.234, containing\_obj=0.727, no\_obj=0.168, cls=0.476

Epoch [20/50], Iter [350/376], Loss: total=2.626, reg=1.243, containing\_obj=0.732, no\_obj=0.169, cls=0.482

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:30<00:00, 41.55it/s]

```
---class aeroplane ap 0.463734176361879---
---class bicycle ap 0.41961030567924384---
---class bird ap 0.3448736285758268---
---class boat ap 0.19848204434841468---
---class bottle ap 0.052481380012674---
---class bus ap 0.4726812076812077---
---class car ap 0.5336402697097424---
---class cat ap 0.521963302512223---
---class chair ap 0.1815518612753533---
---class cow ap 0.2715012236149891---
---class diningtable ap 0.21904761904761905---
---class dog ap 0.36781198029490775---
---class horse ap 0.49332605928915324---
---class motorbike ap 0.47378755714414655---
---class person ap 0.39169918672039705---
---class pottedplant ap 0.12707049419900188---
---class sheep ap 0.10316951972513716---
---class sofa ap 0.2533215701562475---
---class train ap 0.6177718157696752---
---class tvmonitor ap 0.316290172522788---
---map 0.3411907687320313---
19 [0.463734176361879, 0.41961030567924384, 0.3448736285758268, 0.1984820
4434841468, 0.052481380012674, 0.4726812076812077, 0.5336402697097424, 0.
521963302512223, 0.1815518612753533, 0.2715012236149891, 0.21904761904761
905, 0.36781198029490775, 0.49332605928915324, 0.47378755714414655, 0.391
69918672039705, 0.12707049419900188, 0.10316951972513716, 0.2533215701562
475, 0.6177718157696752, 0.316290172522788]
Updating best val loss: 3.15271
```

Starting epoch 21 / 50

Learning Rate for this epoch: 0.001

Epoch [21/50], Iter [50/376], Loss: total=2.365, reg=1.132, containing\_obj=0.692, no\_obj=0.159, cls=0.381

Epoch [21/50], Iter [100/376], Loss: total=2.516, reg=1.206, containing\_obj=0.720, no\_obj=0.167, cls=0.424

Epoch [21/50], Iter [150/376], Loss: total=2.531, reg=1.205, containing\_obj=0.722, no\_obj=0.169, cls=0.436

Epoch [21/50], Iter [200/376], Loss: total=2.538, reg=1.205, containing\_obj=0.717, no\_obj=0.171, cls=0.445

Epoch [21/50], Iter [250/376], Loss: total=2.550, reg=1.211, containing\_obj=0.723, no\_obj=0.171, cls=0.444

Epoch [21/50], Iter [300/376], Loss: total=2.567, reg=1.221, containing\_obj=0.722, no\_obj=0.170, cls=0.454

Epoch [21/50], Iter [350/376], Loss: total=2.574, reg=1.226, containing\_obj=0.717, no\_obj=0.171, cls=0.460

Updating best val loss: 3.11067

Starting epoch 22 / 50

Learning Rate for this epoch: 0.001

Epoch [22/50], Iter [50/376], Loss: total=2.577, reg=1.257, containing\_obj=0.687, no\_obj=0.174, cls=0.458

Epoch [22/50], Iter [100/376], Loss: total=2.533, reg=1.234, containing\_obj=0.682, no\_obj=0.177, cls=0.439

Epoch [22/50], Iter [150/376], Loss: total=2.459, reg=1.187, containing\_obj=0.677, no\_obj=0.179, cls=0.417

Epoch [22/50], Iter [200/376], Loss: total=2.479, reg=1.193, containing\_obj=0.682, no\_obj=0.177, cls=0.427

Epoch [22/50], Iter [250/376], Loss: total=2.493, reg=1.200, containing\_obj=0.687, no\_obj=0.175, cls=0.431

Epoch [22/50], Iter [300/376], Loss: total=2.498, reg=1.201, containing\_obj=0.693, no\_obj=0.175, cls=0.429

Epoch [22/50], Iter [350/376], Loss: total=2.521, reg=1.213, containing\_obj=0.701, no\_obj=0.174, cls=0.433

Starting epoch 23 / 50

Learning Rate for this epoch: 0.001

Epoch [23/50], Iter [50/376], Loss: total=2.365, reg=1.091, containing\_obj=0.700, no\_obj=0.184, cls=0.389

Epoch [23/50], Iter [100/376], Loss: total=2.459, reg=1.179, containing\_obj=0.699, no\_obj=0.182, cls=0.399

Epoch [23/50], Iter [150/376], Loss: total=2.424, reg=1.153, containing\_obj=0.692, no\_obj=0.179, cls=0.400

Epoch [23/50], Iter [200/376], Loss: total=2.452, reg=1.178, containing\_obj=0.700, no\_obj=0.178, cls=0.396

Epoch [23/50], Iter [250/376], Loss: total=2.456, reg=1.176, containing\_obj=0.699, no\_obj=0.177, cls=0.403

Epoch [23/50], Iter [300/376], Loss: total=2.488, reg=1.201, containing\_obj=0.707, no\_obj=0.176, cls=0.404

Epoch [23/50], Iter [350/376], Loss: total=2.486, reg=1.199, containing\_obj=0.708, no\_obj=0.175, cls=0.404

Starting epoch 24 / 50

Learning Rate for this epoch: 0.001

Epoch [24/50], Iter [50/376], Loss: total=2.499, reg=1.209, containing\_obj=0.698, no\_obj=0.173, cls=0.421

Epoch [24/50], Iter [100/376], Loss: total=2.428, reg=1.173, containing\_obj=0.690, no\_obj=0.169, cls=0.396

Epoch [24/50], Iter [150/376], Loss: total=2.432, reg=1.171, containing\_obj=0.688, no\_obj=0.170, cls=0.403

Epoch [24/50], Iter [200/376], Loss: total=2.414, reg=1.167, containing\_obj=0.685, no\_obj=0.171, cls=0.391

Epoch [24/50], Iter [250/376], Loss: total=2.418, reg=1.170, containing\_obj=0.690, no\_obj=0.171, cls=0.387

Epoch [24/50], Iter [300/376], Loss: total=2.449, reg=1.191, containing\_obj=0.694, no\_obj=0.172, cls=0.393

Epoch [24/50], Iter [350/376], Loss: total=2.446, reg=1.187, containing\_obj=0.697, no\_obj=0.173, cls=0.389

Updating best val loss: 3.10698

Starting epoch 25 / 50

Learning Rate for this epoch: 0.001

Epoch [25/50], Iter [50/376], Loss: total=2.414, reg=1.173, containing\_obj=0.681, no\_obj=0.188, cls=0.373

Epoch [25/50], Iter [100/376], Loss: total=2.400, reg=1.175, containing\_obj=0.685, no\_obj=0.183, cls=0.357

Epoch [25/50], Iter [150/376], Loss: total=2.437, reg=1.191, containing\_obj=0.693, no\_obj=0.182, cls=0.370

Epoch [25/50], Iter [200/376], Loss: total=2.418, reg=1.175, containing\_obj=0.692, no\_obj=0.182, cls=0.369

Epoch [25/50], Iter [250/376], Loss: total=2.441, reg=1.179, containing\_obj=0.702, no\_obj=0.181, cls=0.378

Epoch [25/50], Iter [300/376], Loss: total=2.417, reg=1.168, containing\_obj=0.692, no\_obj=0.181, cls=0.376

Epoch [25/50], Iter [350/376], Loss: total=2.410, reg=1.163, containing\_obj=0.688, no\_obj=0.181, cls=0.379

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:29<00:00, 42.31it/s]

```
---class aeroplane ap 0.5218103515577985---
---class bicycle ap 0.4708644592327018---
---class bird ap 0.3294436582210608---
---class boat ap 0.2194489953298088---
---class bottle ap 0.07634967993151967---
---class bus ap 0.429233474248137---
---class car ap 0.5398943111691781---
---class cat ap 0.5520262059729792---
---class chair ap 0.22115497736966336---
---class cow ap 0.3013534906053184---
---class diningtable ap 0.24788971144013158---
---class dog ap 0.4587325670314989---
---class horse ap 0.5623657607662635---
---class motorbike ap 0.38571302992260625---
---class person ap 0.3992836230898037---
---class pottedplant ap 0.09453575983336408---
---class sheep ap 0.2514674463689348---
---class sofa ap 0.3294820069245356---
---class train ap 0.6489954178320607---
---class tvmonitor ap 0.35931971256692563---
---map 0.3699682319707145---
24 [0.5218103515577985, 0.4708644592327018, 0.3294436582210608, 0.2194489
953298088, 0.07634967993151967, 0.429233474248137, 0.5398943111691781, 0.
5520262059729792, 0.22115497736966336, 0.3013534906053184, 0.247889711440
13158, 0.4587325670314989, 0.5623657607662635, 0.38571302992260625, 0.399
2836230898037, 0.09453575983336408, 0.2514674463689348, 0.329482006924535
6, 0.6489954178320607, 0.35931971256692563]
Updating best val loss: 3.09639
```

Starting epoch 26 / 50

Learning Rate for this epoch: 0.001

Epoch [26/50], Iter [50/376], Loss: total=2.409, reg=1.146, containing\_obj=0.716, no\_obj=0.174, cls=0.373

Epoch [26/50], Iter [100/376], Loss: total=2.373, reg=1.120, containing\_obj=0.707, no\_obj=0.179, cls=0.367

Epoch [26/50], Iter [150/376], Loss: total=2.362, reg=1.131, containing\_obj=0.696, no\_obj=0.183, cls=0.351

Epoch [26/50], Iter [200/376], Loss: total=2.331, reg=1.113, containing\_obj=0.689, no\_obj=0.183, cls=0.345

Epoch [26/50], Iter [250/376], Loss: total=2.349, reg=1.120, containing\_obj=0.696, no\_obj=0.183, cls=0.350

Epoch [26/50], Iter [300/376], Loss: total=2.354, reg=1.120, containing\_obj=0.695, no\_obj=0.183, cls=0.355

Epoch [26/50], Iter [350/376], Loss: total=2.366, reg=1.130, containing\_obj=0.701, no\_obj=0.181, cls=0.355

Updating best val loss: 3.07667

Starting epoch 27 / 50

Learning Rate for this epoch: 0.001

Epoch [27/50], Iter [50/376], Loss: total=2.433, reg=1.193, containing\_obj=0.705, no\_obj=0.171, cls=0.363

Epoch [27/50], Iter [100/376], Loss: total=2.381, reg=1.163, containing\_obj=0.684, no\_obj=0.180, cls=0.354

Epoch [27/50], Iter [150/376], Loss: total=2.334, reg=1.130, containing\_obj=0.686, no\_obj=0.182, cls=0.337

Epoch [27/50], Iter [200/376], Loss: total=2.353, reg=1.134, containing\_obj=0.695, no\_obj=0.182, cls=0.342

Epoch [27/50], Iter [250/376], Loss: total=2.349, reg=1.136, containing\_obj=0.688, no\_obj=0.183, cls=0.340

Epoch [27/50], Iter [300/376], Loss: total=2.363, reg=1.146, containing\_obj=0.690, no\_obj=0.183, cls=0.344



Epoch [27/50], Iter [350/376], Loss: total=2.360, reg=1.144, containing\_obj=0.687, no\_obj=0.184, cls=0.345  
Updating best val loss: 3.07555

Starting epoch 28 / 50

Learning Rate for this epoch: 0.001

Epoch [28/50], Iter [50/376], Loss: total=2.030, reg=0.946, containing\_obj=0.603, no\_obj=0.179, cls=0.303  
Epoch [28/50], Iter [100/376], Loss: total=2.123, reg=0.980, containing\_obj=0.631, no\_obj=0.182, cls=0.330  
Epoch [28/50], Iter [150/376], Loss: total=2.264, reg=1.059, containing\_obj=0.675, no\_obj=0.188, cls=0.341  
Epoch [28/50], Iter [200/376], Loss: total=2.247, reg=1.056, containing\_obj=0.660, no\_obj=0.191, cls=0.341  
Epoch [28/50], Iter [250/376], Loss: total=2.257, reg=1.066, containing\_obj=0.665, no\_obj=0.189, cls=0.337  
Epoch [28/50], Iter [300/376], Loss: total=2.259, reg=1.067, containing\_obj=0.676, no\_obj=0.188, cls=0.327  
Epoch [28/50], Iter [350/376], Loss: total=2.272, reg=1.072, containing\_obj=0.679, no\_obj=0.188, cls=0.333  
Updating best val loss: 3.05180

Starting epoch 29 / 50

Learning Rate for this epoch: 0.001

Epoch [29/50], Iter [50/376], Loss: total=2.099, reg=0.987, containing\_obj=0.622, no\_obj=0.192, cls=0.299  
Epoch [29/50], Iter [100/376], Loss: total=2.243, reg=1.074, containing\_obj=0.661, no\_obj=0.186, cls=0.322  
Epoch [29/50], Iter [150/376], Loss: total=2.263, reg=1.083, containing\_obj=0.675, no\_obj=0.188, cls=0.317  
Epoch [29/50], Iter [200/376], Loss: total=2.238, reg=1.074, containing\_obj=0.668, no\_obj=0.185, cls=0.311  
Epoch [29/50], Iter [250/376], Loss: total=2.240, reg=1.075, containing\_obj=0.668, no\_obj=0.187, cls=0.311  
Epoch [29/50], Iter [300/376], Loss: total=2.260, reg=1.080, containing\_obj=0.672, no\_obj=0.187, cls=0.321  
Epoch [29/50], Iter [350/376], Loss: total=2.264, reg=1.077, containing\_obj=0.672, no\_obj=0.188, cls=0.327

Starting epoch 30 / 50

Learning Rate for this epoch: 0.001

Epoch [30/50], Iter [50/376], Loss: total=2.315, reg=1.103, containing\_obj=0.737, no\_obj=0.187, cls=0.288  
Epoch [30/50], Iter [100/376], Loss: total=2.241, reg=1.073, containing\_obj=0.680, no\_obj=0.187, cls=0.301  
Epoch [30/50], Iter [150/376], Loss: total=2.277, reg=1.074, containing\_obj=0.694, no\_obj=0.185, cls=0.324  
Epoch [30/50], Iter [200/376], Loss: total=2.258, reg=1.065, containing\_obj=0.683, no\_obj=0.188, cls=0.321  
Epoch [30/50], Iter [250/376], Loss: total=2.246, reg=1.059, containing\_obj=0.675, no\_obj=0.188, cls=0.324  
Epoch [30/50], Iter [300/376], Loss: total=2.229, reg=1.052, containing\_obj=0.673, no\_obj=0.188, cls=0.315  
Epoch [30/50], Iter [350/376], Loss: total=2.231, reg=1.055, containing\_obj=0.675, no\_obj=0.187, cls=0.314

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:30<00:00, 41.64it/s]

```
---class aeroplane ap 0.5467547067562715---
---class bicycle ap 0.2863469863469863---
---class bird ap 0.4170628429507444---
---class boat ap 0.24993316829057033---
---class bottle ap 0.09088491633972368---
---class bus ap 0.5725843951466496---
---class car ap 0.5278500722928762---
---class cat ap 0.5877564306385779---
---class chair ap 0.1460357809678931---
---class cow ap 0.27852150965755457---
---class diningtable ap 0.285407496331866---
---class dog ap 0.5024480837649319---
---class horse ap 0.5257528702147107---
---class motorbike ap 0.5020982218475488---
---class person ap 0.42966190804762466---
---class pottedplant ap 0.13467542955281012---
---class sheep ap 0.2370838302764299---
---class sofa ap 0.34880121568094735---
---class train ap 0.6826296942612603---
---class tvmonitor ap 0.4977820749551855---
---map 0.39250358171605815---
29 [0.5467547067562715, 0.2863469863469863, 0.4170628429507444, 0.2499331
6829057033, 0.09088491633972368, 0.5725843951466496, 0.5278500722928762,
0.5877564306385779, 0.1460357809678931, 0.27852150965755457, 0.2854074963
31866, 0.5024480837649319, 0.5257528702147107, 0.5020982218475488, 0.4296
6190804762466, 0.13467542955281012, 0.2370838302764299, 0.348801215680947
35, 0.6826296942612603, 0.4977820749551855]
```

Starting epoch 31 / 50

Learning Rate for this epoch: 0.0001

Epoch [31/50], Iter [50/376], Loss: total=2.122, reg=0.992, containing\_obj=0.632, no\_obj=0.205, cls=0.293

Epoch [31/50], Iter [100/376], Loss: total=2.153, reg=1.027, containing\_obj=0.643, no\_obj=0.202, cls=0.282

Epoch [31/50], Iter [150/376], Loss: total=2.093, reg=0.992, containing\_obj=0.636, no\_obj=0.200, cls=0.265

Epoch [31/50], Iter [200/376], Loss: total=2.081, reg=0.987, containing\_obj=0.634, no\_obj=0.199, cls=0.262

Epoch [31/50], Iter [250/376], Loss: total=2.096, reg=0.991, containing\_obj=0.632, no\_obj=0.199, cls=0.274

Epoch [31/50], Iter [300/376], Loss: total=2.098, reg=0.988, containing\_obj=0.636, no\_obj=0.199, cls=0.275

Epoch [31/50], Iter [350/376], Loss: total=2.083, reg=0.981, containing\_obj=0.634, no\_obj=0.199, cls=0.268

Updating best val loss: 2.91346

Starting epoch 32 / 50

Learning Rate for this epoch: 0.0001

Epoch [32/50], Iter [50/376], Loss: total=2.146, reg=1.013, containing\_obj=0.676, no\_obj=0.194, cls=0.263

Epoch [32/50], Iter [100/376], Loss: total=2.078, reg=0.969, containing\_obj=0.655, no\_obj=0.197, cls=0.257

Epoch [32/50], Iter [150/376], Loss: total=2.008, reg=0.938, containing\_obj=0.621, no\_obj=0.198, cls=0.251

Epoch [32/50], Iter [200/376], Loss: total=2.024, reg=0.956, containing\_obj=0.629, no\_obj=0.197, cls=0.243

Epoch [32/50], Iter [250/376], Loss: total=2.049, reg=0.967, containing\_obj=0.639, no\_obj=0.195, cls=0.248

Epoch [32/50], Iter [300/376], Loss: total=2.023, reg=0.955, containing\_obj=0.629, no\_obj=0.196, cls=0.243

Epoch [32/50], Iter [350/376], Loss: total=2.021, reg=0.955, containing\_obj=

bj=0.630, no\_obj=0.196, cls=0.240  
Updating best val loss: 2.89792

Starting epoch 33 / 50

Learning Rate for this epoch: 0.0001

Epoch [33/50], Iter [50/376], Loss: total=1.910, reg=0.911, containing\_obj=0.598, no\_obj=0.194, cls=0.208

Epoch [33/50], Iter [100/376], Loss: total=1.934, reg=0.910, containing\_obj=0.603, no\_obj=0.197, cls=0.224

Epoch [33/50], Iter [150/376], Loss: total=1.952, reg=0.928, containing\_obj=0.602, no\_obj=0.194, cls=0.227

Epoch [33/50], Iter [200/376], Loss: total=1.958, reg=0.924, containing\_obj=0.609, no\_obj=0.193, cls=0.232

Epoch [33/50], Iter [250/376], Loss: total=1.960, reg=0.931, containing\_obj=0.607, no\_obj=0.193, cls=0.228

Epoch [33/50], Iter [300/376], Loss: total=1.975, reg=0.937, containing\_obj=0.616, no\_obj=0.193, cls=0.229

Epoch [33/50], Iter [350/376], Loss: total=1.978, reg=0.937, containing\_obj=0.618, no\_obj=0.197, cls=0.226

Updating best val loss: 2.89615

Starting epoch 34 / 50

Learning Rate for this epoch: 0.0001

Epoch [34/50], Iter [50/376], Loss: total=1.874, reg=0.876, containing\_obj=0.580, no\_obj=0.201, cls=0.217

Epoch [34/50], Iter [100/376], Loss: total=1.961, reg=0.910, containing\_obj=0.633, no\_obj=0.199, cls=0.219

Epoch [34/50], Iter [150/376], Loss: total=1.978, reg=0.929, containing\_obj=0.630, no\_obj=0.200, cls=0.219

Epoch [34/50], Iter [200/376], Loss: total=1.988, reg=0.936, containing\_obj=0.628, no\_obj=0.199, cls=0.223

Epoch [34/50], Iter [250/376], Loss: total=1.972, reg=0.926, containing\_obj=0.625, no\_obj=0.199, cls=0.222

Epoch [34/50], Iter [300/376], Loss: total=1.965, reg=0.922, containing\_obj=0.624, no\_obj=0.198, cls=0.222

Epoch [34/50], Iter [350/376], Loss: total=1.963, reg=0.918, containing\_obj=0.624, no\_obj=0.196, cls=0.224

Starting epoch 35 / 50

Learning Rate for this epoch: 0.0001

Epoch [35/50], Iter [50/376], Loss: total=1.883, reg=0.898, containing\_obj=0.602, no\_obj=0.206, cls=0.176

Epoch [35/50], Iter [100/376], Loss: total=1.934, reg=0.922, containing\_obj=0.615, no\_obj=0.198, cls=0.199

Epoch [35/50], Iter [150/376], Loss: total=1.930, reg=0.924, containing\_obj=0.602, no\_obj=0.196, cls=0.208

Epoch [35/50], Iter [200/376], Loss: total=1.953, reg=0.933, containing\_obj=0.617, no\_obj=0.194, cls=0.209

Epoch [35/50], Iter [250/376], Loss: total=1.974, reg=0.946, containing\_obj=0.616, no\_obj=0.196, cls=0.217

Epoch [35/50], Iter [300/376], Loss: total=1.984, reg=0.952, containing\_obj=0.623, no\_obj=0.195, cls=0.213

Epoch [35/50], Iter [350/376], Loss: total=1.974, reg=0.948, containing\_obj=0.619, no\_obj=0.195, cls=0.212

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:30<00:00, 41.57it/s]

```
---class aeroplane ap 0.6377968686238712---
---class bicycle ap 0.5667887280495053---
---class bird ap 0.46997222536825306---
---class boat ap 0.37182674467463855---
---class bottle ap 0.09680341688317132---
---class bus ap 0.5790651220639974---
---class car ap 0.6075175676983549---
---class cat ap 0.6120606141089372---
---class chair ap 0.24257145866370386---
---class cow ap 0.44134634392027794---
---class diningtable ap 0.4123018004078838---
---class dog ap 0.5482859765218625---
---class horse ap 0.6060001096126354---
---class motorbike ap 0.4835638087882555---
---class person ap 0.4823366068407742---
---class pottedplant ap 0.16793959773763412---
---class sheep ap 0.21527094867099883---
---class sofa ap 0.44712695289110366---
---class train ap 0.7427171182784261---
---class tvmonitor ap 0.4859033888937304---
---map 0.46085976993490074---
34 [0.6377968686238712, 0.5667887280495053, 0.46997222536825306, 0.371826
74467463855, 0.09680341688317132, 0.5790651220639974, 0.6075175676983549,
0.6120606141089372, 0.24257145866370386, 0.44134634392027794, 0.412301800
4078838, 0.5482859765218625, 0.6060001096126354, 0.4835638087882555, 0.48
23366068407742, 0.16793959773763412, 0.21527094867099883, 0.4471269528911
0366, 0.7427171182784261, 0.4859033888937304]
Updating best val loss: 2.88914
```

Starting epoch 36 / 50

Learning Rate for this epoch: 0.0001

Epoch [36/50], Iter [50/376], Loss: total=1.925, reg=0.911, containing\_obj=0.615, no\_obj=0.189, cls=0.209

Epoch [36/50], Iter [100/376], Loss: total=1.920, reg=0.908, containing\_obj=0.616, no\_obj=0.191, cls=0.204

Epoch [36/50], Iter [150/376], Loss: total=1.945, reg=0.919, containing\_obj=0.622, no\_obj=0.194, cls=0.210

Epoch [36/50], Iter [200/376], Loss: total=1.932, reg=0.908, containing\_obj=0.615, no\_obj=0.196, cls=0.213

Epoch [36/50], Iter [250/376], Loss: total=1.925, reg=0.906, containing\_obj=0.610, no\_obj=0.196, cls=0.212

Epoch [36/50], Iter [300/376], Loss: total=1.933, reg=0.912, containing\_obj=0.610, no\_obj=0.198, cls=0.213

Epoch [36/50], Iter [350/376], Loss: total=1.932, reg=0.909, containing\_obj=0.613, no\_obj=0.199, cls=0.211

Updating best val loss: 2.88862

Starting epoch 37 / 50

Learning Rate for this epoch: 0.0001

Epoch [37/50], Iter [50/376], Loss: total=1.929, reg=0.904, containing\_obj=0.611, no\_obj=0.188, cls=0.226

Epoch [37/50], Iter [100/376], Loss: total=1.914, reg=0.899, containing\_obj=0.613, no\_obj=0.189, cls=0.212

Epoch [37/50], Iter [150/376], Loss: total=1.915, reg=0.907, containing\_obj=0.602, no\_obj=0.192, cls=0.214

Epoch [37/50], Iter [200/376], Loss: total=1.897, reg=0.891, containing\_obj=0.601, no\_obj=0.193, cls=0.212

Epoch [37/50], Iter [250/376], Loss: total=1.893, reg=0.882, containing\_obj=0.608, no\_obj=0.195, cls=0.208

Epoch [37/50], Iter [300/376], Loss: total=1.883, reg=0.882, containing\_obj=0.602, no\_obj=0.196, cls=0.203

Epoch [37/50], Iter [350/376], Loss: total=1.890, reg=0.892, containing\_obj=0.600, no\_obj=0.196, cls=0.201

Starting epoch 38 / 50

Learning Rate for this epoch: 0.0001

Epoch [38/50], Iter [50/376], Loss: total=2.028, reg=0.960, containing\_obj=0.655, no\_obj=0.197, cls=0.216

Epoch [38/50], Iter [100/376], Loss: total=1.961, reg=0.925, containing\_obj=0.614, no\_obj=0.203, cls=0.219

Epoch [38/50], Iter [150/376], Loss: total=1.926, reg=0.905, containing\_obj=0.602, no\_obj=0.201, cls=0.219

Epoch [38/50], Iter [200/376], Loss: total=1.889, reg=0.881, containing\_obj=0.599, no\_obj=0.197, cls=0.212

Epoch [38/50], Iter [250/376], Loss: total=1.888, reg=0.887, containing\_obj=0.602, no\_obj=0.195, cls=0.204

Epoch [38/50], Iter [300/376], Loss: total=1.893, reg=0.888, containing\_obj=0.600, no\_obj=0.197, cls=0.207

Epoch [38/50], Iter [350/376], Loss: total=1.904, reg=0.894, containing\_obj=0.606, no\_obj=0.196, cls=0.207

Updating best val loss: 2.87321

Starting epoch 39 / 50

Learning Rate for this epoch: 0.0001

Epoch [39/50], Iter [50/376], Loss: total=1.810, reg=0.840, containing\_obj=0.574, no\_obj=0.206, cls=0.190

Epoch [39/50], Iter [100/376], Loss: total=1.797, reg=0.850, containing\_obj=0.569, no\_obj=0.197, cls=0.181

Epoch [39/50], Iter [150/376], Loss: total=1.807, reg=0.854, containing\_obj=0.565, no\_obj=0.200, cls=0.188

Epoch [39/50], Iter [200/376], Loss: total=1.860, reg=0.876, containing\_obj=0.590, no\_obj=0.197, cls=0.197

Epoch [39/50], Iter [250/376], Loss: total=1.893, reg=0.889, containing\_obj=0.603, no\_obj=0.200, cls=0.201

Epoch [39/50], Iter [300/376], Loss: total=1.883, reg=0.884, containing\_obj=0.599, no\_obj=0.200, cls=0.200

Epoch [39/50], Iter [350/376], Loss: total=1.878, reg=0.878, containing\_obj=0.600, no\_obj=0.202, cls=0.199

Updating best val loss: 2.86539

Starting epoch 40 / 50

Learning Rate for this epoch: 0.0001

Epoch [40/50], Iter [50/376], Loss: total=1.770, reg=0.829, containing\_obj=0.563, no\_obj=0.199, cls=0.178

Epoch [40/50], Iter [100/376], Loss: total=1.897, reg=0.905, containing\_obj=0.605, no\_obj=0.197, cls=0.190

Epoch [40/50], Iter [150/376], Loss: total=1.908, reg=0.912, containing\_obj=0.608, no\_obj=0.198, cls=0.191

Epoch [40/50], Iter [200/376], Loss: total=1.887, reg=0.904, containing\_obj=0.597, no\_obj=0.198, cls=0.188

Epoch [40/50], Iter [250/376], Loss: total=1.918, reg=0.920, containing\_obj=0.601, no\_obj=0.199, cls=0.199

Epoch [40/50], Iter [300/376], Loss: total=1.917, reg=0.919, containing\_obj=0.601, no\_obj=0.198, cls=0.198

Epoch [40/50], Iter [350/376], Loss: total=1.900, reg=0.907, containing\_obj=0.594, no\_obj=0.199, cls=0.200

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:31<00:00, 40.46it/s]

```
---class aeroplane ap 0.6066087445192966---
---class bicycle ap 0.6281836594807055---
---class bird ap 0.48134857753185806---
---class boat ap 0.33467068117301446---
---class bottle ap 0.11195236760820834---
---class bus ap 0.6042953252921691---
---class car ap 0.5999580808788909---
---class cat ap 0.6784794915498209---
---class chair ap 0.27070023683480604---
---class cow ap 0.41359863185121226---
---class diningtable ap 0.3970850438792401---
---class dog ap 0.5000666930655675---
---class horse ap 0.6227635489522279---
---class motorbike ap 0.5057718706736282---
---class person ap 0.4811362308326159---
---class pottedplant ap 0.2084698489485198---
---class sheep ap 0.23708982197908138---
---class sofa ap 0.4653092890264147---
---class train ap 0.6976349222635745---
---class tvmonitor ap 0.5007717640738744---
---map 0.46729474152073636---
39 [0.6066087445192966, 0.6281836594807055, 0.48134857753185806, 0.334670
68117301446, 0.11195236760820834, 0.6042953252921691, 0.5999580808788909,
0.6784794915498209, 0.27070023683480604, 0.41359863185121226, 0.397085043
8792401, 0.5000666930655675, 0.6227635489522279, 0.5057718706736282, 0.48
11362308326159, 0.2084698489485198, 0.23708982197908138, 0.46530928902641
47, 0.6976349222635745, 0.5007717640738744]
Updating best val loss: 2.86508
```

Starting epoch 41 / 50

Learning Rate for this epoch: 1e-05

Epoch [41/50], Iter [50/376], Loss: total=1.921, reg=0.917, containing\_obj=0.605, no\_obj=0.201, cls=0.197

Epoch [41/50], Iter [100/376], Loss: total=1.913, reg=0.923, containing\_obj=0.598, no\_obj=0.199, cls=0.194

Epoch [41/50], Iter [150/376], Loss: total=1.889, reg=0.901, containing\_obj=0.599, no\_obj=0.197, cls=0.192

Epoch [41/50], Iter [200/376], Loss: total=1.885, reg=0.893, containing\_obj=0.596, no\_obj=0.196, cls=0.199

Epoch [41/50], Iter [250/376], Loss: total=1.910, reg=0.909, containing\_obj=0.602, no\_obj=0.197, cls=0.203

Epoch [41/50], Iter [300/376], Loss: total=1.893, reg=0.901, containing\_obj=0.598, no\_obj=0.197, cls=0.197

Epoch [41/50], Iter [350/376], Loss: total=1.876, reg=0.886, containing\_obj=0.598, no\_obj=0.197, cls=0.194

Updating best val loss: 2.85467

Starting epoch 42 / 50

Learning Rate for this epoch: 1e-05

Epoch [42/50], Iter [50/376], Loss: total=1.901, reg=0.853, containing\_obj=0.660, no\_obj=0.198, cls=0.190

Epoch [42/50], Iter [100/376], Loss: total=1.934, reg=0.885, containing\_obj=0.645, no\_obj=0.205, cls=0.199

Epoch [42/50], Iter [150/376], Loss: total=1.928, reg=0.887, containing\_obj=0.640, no\_obj=0.201, cls=0.200

Epoch [42/50], Iter [200/376], Loss: total=1.891, reg=0.874, containing\_obj=0.621, no\_obj=0.199, cls=0.197

Epoch [42/50], Iter [250/376], Loss: total=1.894, reg=0.873, containing\_obj=0.625, no\_obj=0.198, cls=0.198

Epoch [42/50], Iter [300/376], Loss: total=1.868, reg=0.861, containing\_obj=0.610, no\_obj=0.198, cls=0.199

Epoch [42/50], Iter [350/376], Loss: total=1.853, reg=0.857, containing\_obj=0.604, no\_obj=0.199, cls=0.193

Starting epoch 43 / 50

Learning Rate for this epoch: 1e-05

Epoch [43/50], Iter [50/376], Loss: total=1.805, reg=0.835, containing\_obj=0.582, no\_obj=0.198, cls=0.189

Epoch [43/50], Iter [100/376], Loss: total=1.842, reg=0.862, containing\_obj=0.592, no\_obj=0.203, cls=0.185

Epoch [43/50], Iter [150/376], Loss: total=1.882, reg=0.875, containing\_obj=0.616, no\_obj=0.200, cls=0.192

Epoch [43/50], Iter [200/376], Loss: total=1.867, reg=0.868, containing\_obj=0.607, no\_obj=0.200, cls=0.193

Epoch [43/50], Iter [250/376], Loss: total=1.866, reg=0.865, containing\_obj=0.606, no\_obj=0.200, cls=0.194

Epoch [43/50], Iter [300/376], Loss: total=1.855, reg=0.865, containing\_obj=0.597, no\_obj=0.201, cls=0.191

Epoch [43/50], Iter [350/376], Loss: total=1.862, reg=0.866, containing\_obj=0.601, no\_obj=0.202, cls=0.192

Starting epoch 44 / 50

Learning Rate for this epoch: 1e-05

Epoch [44/50], Iter [50/376], Loss: total=2.012, reg=0.918, containing\_obj=0.669, no\_obj=0.195, cls=0.230

Epoch [44/50], Iter [100/376], Loss: total=1.926, reg=0.881, containing\_obj=0.635, no\_obj=0.199, cls=0.210

Epoch [44/50], Iter [150/376], Loss: total=1.882, reg=0.864, containing\_obj=0.618, no\_obj=0.196, cls=0.203

Epoch [44/50], Iter [200/376], Loss: total=1.877, reg=0.873, containing\_obj=0.612, no\_obj=0.195, cls=0.197

Epoch [44/50], Iter [250/376], Loss: total=1.877, reg=0.874, containing\_obj=0.611, no\_obj=0.197, cls=0.195

Epoch [44/50], Iter [300/376], Loss: total=1.867, reg=0.872, containing\_obj=0.609, no\_obj=0.197, cls=0.190

Epoch [44/50], Iter [350/376], Loss: total=1.855, reg=0.867, containing\_obj=0.605, no\_obj=0.196, cls=0.187

Starting epoch 45 / 50

Learning Rate for this epoch: 1e-05

Epoch [45/50], Iter [50/376], Loss: total=1.874, reg=0.894, containing\_obj=0.598, no\_obj=0.200, cls=0.182

Epoch [45/50], Iter [100/376], Loss: total=1.856, reg=0.890, containing\_obj=0.590, no\_obj=0.199, cls=0.177

Epoch [45/50], Iter [150/376], Loss: total=1.943, reg=0.932, containing\_obj=0.613, no\_obj=0.200, cls=0.198

Epoch [45/50], Iter [200/376], Loss: total=1.884, reg=0.896, containing\_obj=0.595, no\_obj=0.199, cls=0.194

Epoch [45/50], Iter [250/376], Loss: total=1.869, reg=0.885, containing\_obj=0.594, no\_obj=0.198, cls=0.192

Epoch [45/50], Iter [300/376], Loss: total=1.872, reg=0.884, containing\_obj=0.595, no\_obj=0.198, cls=0.195

Epoch [45/50], Iter [350/376], Loss: total=1.876, reg=0.886, containing\_obj=0.598, no\_obj=0.198, cls=0.193

---Evaluate model on test samples---

100%|██████████| 1255/1255 [00:31<00:00, 39.68it/s]

```
---class aeroplane ap 0.6315298082012735---
---class bicycle ap 0.6063405225886289---
---class bird ap 0.4748701324428052---
---class boat ap 0.3379636587710013---
---class bottle ap 0.12907940871832974---
---class bus ap 0.58675829597474---
---class car ap 0.5952232554253093---
---class cat ap 0.6942820852814685---
---class chair ap 0.2875845859646111---
---class cow ap 0.4112014098300425---
---class diningtable ap 0.39412874338872167---
---class dog ap 0.5408598337314102---
---class horse ap 0.6077761749065336---
---class motorbike ap 0.4971991777340232---
---class person ap 0.4781225518822114---
---class pottedplant ap 0.16282772016168517---
---class sheep ap 0.23349715527705323---
---class sofa ap 0.4777071774878793---
---class train ap 0.7342102167256553---
---class tvmonitor ap 0.5013883516527632---
---map 0.4691275133073073---
44 [0.6315298082012735, 0.6063405225886289, 0.4748701324428052, 0.3379636
587710013, 0.12907940871832974, 0.58675829597474, 0.5952232554253093, 0.6
942820852814685, 0.2875845859646111, 0.4112014098300425, 0.39412874338872
167, 0.5408598337314102, 0.6077761749065336, 0.4971991777340232, 0.478122
5518822114, 0.16282772016168517, 0.23349715527705323, 0.4777071774878793,
0.7342102167256553, 0.5013883516527632]
```



Starting epoch 46 / 50

Learning Rate for this epoch: 1e-05

Epoch [46/50], Iter [50/376], Loss: total=1.770, reg=0.793, containing\_obj=0.581, no\_obj=0.200, cls=0.196

Epoch [46/50], Iter [100/376], Loss: total=1.835, reg=0.838, containing\_obj=0.594, no\_obj=0.202, cls=0.202

Epoch [46/50], Iter [150/376], Loss: total=1.898, reg=0.880, containing\_obj=0.621, no\_obj=0.198, cls=0.199

Epoch [46/50], Iter [200/376], Loss: total=1.880, reg=0.880, containing\_obj=0.608, no\_obj=0.197, cls=0.195

Epoch [46/50], Iter [250/376], Loss: total=1.837, reg=0.852, containing\_obj=0.597, no\_obj=0.200, cls=0.188

Epoch [46/50], Iter [300/376], Loss: total=1.823, reg=0.842, containing\_obj=0.592, no\_obj=0.199, cls=0.190

Epoch [46/50], Iter [350/376], Loss: total=1.825, reg=0.844, containing\_obj=0.592, no\_obj=0.198, cls=0.191

Starting epoch 47 / 50

Learning Rate for this epoch: 1e-05

Epoch [47/50], Iter [50/376], Loss: total=1.801, reg=0.813, containing\_obj=0.587, no\_obj=0.198, cls=0.203

Epoch [47/50], Iter [100/376], Loss: total=1.876, reg=0.868, containing\_obj=0.618, no\_obj=0.197, cls=0.193

Epoch [47/50], Iter [150/376], Loss: total=1.856, reg=0.858, containing\_obj=0.606, no\_obj=0.198, cls=0.193

Epoch [47/50], Iter [200/376], Loss: total=1.822, reg=0.847, containing\_obj=0.594, no\_obj=0.196, cls=0.184

Epoch [47/50], Iter [250/376], Loss: total=1.812, reg=0.842, containing\_obj=0.591, no\_obj=0.197, cls=0.182

Epoch [47/50], Iter [300/376], Loss: total=1.812, reg=0.839, containing\_obj=0.596, no\_obj=0.197, cls=0.181

Epoch [47/50], Iter [350/376], Loss: total=1.826, reg=0.848, containing\_obj=0.596, no\_obj=0.198, cls=0.184

Starting epoch 48 / 50

Learning Rate for this epoch: 1e-05

Epoch [48/50], Iter [50/376], Loss: total=1.882, reg=0.886, containing\_obj=0.598, no\_obj=0.192, cls=0.206

Epoch [48/50], Iter [100/376], Loss: total=1.858, reg=0.874, containing\_obj=0.592, no\_obj=0.192, cls=0.200

Epoch [48/50], Iter [150/376], Loss: total=1.857, reg=0.871, containing\_obj=0.598, no\_obj=0.195, cls=0.193

Epoch [48/50], Iter [200/376], Loss: total=1.871, reg=0.883, containing\_obj=0.598, no\_obj=0.196, cls=0.193

Epoch [48/50], Iter [250/376], Loss: total=1.850, reg=0.868, containing\_obj=0.595, no\_obj=0.198, cls=0.189

Epoch [48/50], Iter [300/376], Loss: total=1.849, reg=0.868, containing\_obj=0.596, no\_obj=0.198, cls=0.188

Epoch [48/50], Iter [350/376], Loss: total=1.852, reg=0.868, containing\_obj=0.600, no\_obj=0.198, cls=0.185

Starting epoch 49 / 50

Learning Rate for this epoch: 1e-05

Epoch [49/50], Iter [50/376], Loss: total=1.855, reg=0.875, containing\_obj=0.600, no\_obj=0.201, cls=0.180

Epoch [49/50], Iter [100/376], Loss: total=1.854, reg=0.873, containing\_obj=0.605, no\_obj=0.201, cls=0.175

Epoch [49/50], Iter [150/376], Loss: total=1.828, reg=0.855, containing\_obj=0.599, no\_obj=0.200, cls=0.174

```
Epoch [49/50], Iter [200/376], Loss: total=1.796, reg=0.841, containing_o
bj=0.585, no_obj=0.199, cls=0.171
Epoch [49/50], Iter [250/376], Loss: total=1.796, reg=0.845, containing_o
bj=0.583, no_obj=0.199, cls=0.169
Epoch [49/50], Iter [300/376], Loss: total=1.801, reg=0.845, containing_o
bj=0.587, no_obj=0.199, cls=0.170
Epoch [49/50], Iter [350/376], Loss: total=1.814, reg=0.845, containing_o
bj=0.594, no_obj=0.199, cls=0.176
Updating best val loss: 2.85190
```

Starting epoch 50 / 50

Learning Rate for this epoch: 1e-05

```
Epoch [50/50], Iter [50/376], Loss: total=1.835, reg=0.874, containing_ob
j=0.590, no_obj=0.200, cls=0.171
Epoch [50/50], Iter [100/376], Loss: total=1.887, reg=0.887, containing_o
bj=0.616, no_obj=0.201, cls=0.183
Epoch [50/50], Iter [150/376], Loss: total=1.880, reg=0.880, containing_o
bj=0.613, no_obj=0.201, cls=0.186
Epoch [50/50], Iter [200/376], Loss: total=1.880, reg=0.880, containing_o
bj=0.609, no_obj=0.203, cls=0.188
Epoch [50/50], Iter [250/376], Loss: total=1.850, reg=0.860, containing_o
bj=0.598, no_obj=0.202, cls=0.191
Epoch [50/50], Iter [300/376], Loss: total=1.858, reg=0.871, containing_o
bj=0.599, no_obj=0.200, cls=0.188
Epoch [50/50], Iter [350/376], Loss: total=1.859, reg=0.872, containing_o
bj=0.598, no_obj=0.199, cls=0.190
---Evaluate model on test samples---
```

100%|██████████| 1255/1255 [00:31<00:00, 39.32it/s]

```
---class aeroplane ap 0.615949735601417---
---class bicycle ap 0.6157775321704524---
---class bird ap 0.4810726534511093---
---class boat ap 0.3178122370921667---
---class bottle ap 0.1078000151010163---
---class bus ap 0.5471419164722934---
---class car ap 0.5769939332811684---
---class cat ap 0.697558629482799---
---class chair ap 0.28496726010042256---
---class cow ap 0.43116441550961815---
---class diningtable ap 0.39881556890649394---
---class dog ap 0.5418315921532685---
---class horse ap 0.6021238639196907---
---class motorbike ap 0.5365404275951765---
---class person ap 0.45401377308644064---
---class pottedplant ap 0.15835678113767676---
---class sheep ap 0.22214626489230463---
---class sofa ap 0.4785335245351405---
---class train ap 0.7137208695164142---
---class tvmonitor ap 0.5034555538904202---
---map 0.46428882739477456---
```

```
49 [0.615949735601417, 0.6157775321704524, 0.4810726534511093, 0.31781223
70921667, 0.1078000151010163, 0.5471419164722934, 0.5769939332811684, 0.6
97558629482799, 0.28496726010042256, 0.43116441550961815, 0.3988155689064
9394, 0.5418315921532685, 0.6021238639196907, 0.5365404275951765, 0.45401
377308644064, 0.15835678113767676, 0.22214626489230463, 0.478533524535140
5, 0.7137208695164142, 0.5034555538904202]
```

## View example predictions


```
In [ ]: ckpt = torch.load('checkpoints/best_detector.pth')
net = resnet50()
net.load_state_dict(ckpt)
net.to(device)
net.eval()

# select random image from val set
image_name = random.choice(val_dataset.fnames)
image = cv2.imread(os.path.join(file_root_val, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_val)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX)
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
        color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.FONT_HERSHEY_

plt.figure(figsize = (15,15))
plt.imshow(image)
```

```
predicting...
Out[ ]: <matplotlib.image.AxesImage at 0x7f7b7019b450>
```



Kaggle submission (85%)

## Predict Result

Predict the results based on testing set. Upload to [Kaggle](#).

### How to upload

1. Click the folder icon in the left hand side of Colab.
2. Right click "result.csv". Select "Download"
3. To kaggle. Click "Submit Predictions"
4. Upload the result.csv
5. System will automaticlaly calculate the accuracy of 50% dataset and publish this result to leaderboard.

---

預測 test 並將結果上傳至Kaggle。 [連結](#)

執行完畢此區的程式碼後，會將 test 預測完的結果存下來。

### 上傳流程

1. 點選左側選單最下方的資料夾圖示
2. 右鍵「result.csv」
3. 點選「Download」
4. 至連結網頁點選「Submit Predictions」
5. 將剛剛下載的檔案上傳
6. 系統會計算並公布其中50%資料的正確率

```
In [16]: root_test = 'data/VOCdevkit_2007/VOC2007test/JPEGImages/'
         file_test = 'data/voc2007test.txt'

         ckpt = torch.load('checkpoints/best_detector.pth')
         net = resnet50()
         net.load_state_dict(ckpt)
         net = net.to(device)
```

By using the test\_evaluate function, you will obtain predictions for each image.

```
In [17]: preds_submission = test_evaluate(net, test_dataset_file=file_test, img_ro

---Evaluate model on test samples---
 0%|          | 0/4950 [00:00<?, ?it/s]100%|██████████| 4950/4950 [02:03
<00:00, 40.12it/s]
```

The write\_csv function will use preds\_submission to write into a CSV file called 'result.csv'.

```
In [18]: write_csv(preds_submission)
```

## Report (15%)

In your report, please include:

- a. A brief discussion on your implementation.
- b. Report the best train and validation accuracy in all of your experiments and discuss any strategies or tricks you've employed.
- c. Report the results for extra credits and also provide a discussion, if any.

## Extra Credit (15%)

- Pick a fun video like [this one](#), run your detector on it (a subset of frames would be OK), and produce a video showing your results.
- Try to replace the provided pre-trained network with a different one and train with the YOLO loss on top to attempt to get better accuracy.
- Or any other methods that you try to improve the performance.

```

In [ ]: import cv2
import torch
import torchvision.transforms as transforms
from src.config import YOLO_IMG_DIM, VOC_IMG_MEAN, VOC_CLASSES, COLORS
from torch.autograd import Variable
from src.predict import decoder

net = resnet50(pretrained=False)
net.load_state_dict(torch.load('checkpoints/detector.pth'))
net.eval()
net.to(device)

# 视频读取
cap = cv2.VideoCapture('snl-digital-short-yolo-snl.mp4')
fourcc = cv2.VideoWriter_fourcc(*'mp4v')

# 获取视频的宽度和高度
frame_width = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
frame_height = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))

out = cv2.VideoWriter('snl-digital-short-yolo-snl_output.mp4', fourcc, 20

while cap.isOpened():
    ret, image = cap.read() # ret, frame
    if not ret:
        break

    result = []
    frame = cv2.resize(image, (YOLO_IMG_DIM, YOLO_IMG_DIM)) # YOLO_IMG_DIM
    frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)

    #img = cv2.resize(image, (YOLO_IMG_DIM, YOLO_IMG_DIM))
    h, w, _ = frame.shape
    img = cv2.resize(image, (YOLO_IMG_DIM, YOLO_IMG_DIM)) # YOLO_IMG_DIM
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    mean = VOC_IMG_MEAN
    img = img - np.array(mean, dtype=np.float32)

    transform = transforms.Compose(
        [
            transforms.ToTensor(),
        ]
    )
    img = transform(img)
    with torch.no_grad():
        img = Variable(img[None, :, :, :])
        img = img.cuda()

        pred = net(img) # 1xSxSx(B*5+C)
        pred = pred.cpu()
        boxes, cls_indexs, probs = decoder(pred)

        for i, box in enumerate(boxes):
            x1 = int(box[0] * w)
            x2 = int(box[2] * w)
            y1 = int(box[1] * h)
            y2 = int(box[3] * h)
            cls_index = cls_indexs[i]
            cls_index = int(cls_index) # convert LongTensor to int
            prob = probs[i]
            prob = float(prob)

```

```

        prob = float(prob)
        result.append(
            [(x1, y1), (x2, y2), VOC_CLASSES[cls_index], prob]
        )
    for left_up, right_bottom, class_name, prob in result:
        color = COLORS[VOC_CLASSES.index(class_name)]
        cv2.rectangle(frame, left_up, right_bottom, color, 2)
        label = class_name + str(round(prob, 2))
        text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_
        p1 = (left_up[0], left_up[1] - text_size[1])
        cv2.rectangle(frame, (p1[0] - 2 // 2, p1[1] - 2 - baseline),
                        color, -1)
        cv2.putText(frame, label, (p1[0], p1[1] + baseline), cv2.FONT

    out.write(frame)

cap.release()
out.release()

```

The video can be seen through the link below:

<https://drive.google.com/file/d/1lc5o3AT2AwelwwkCqVZHt2gqJPFCELQT/view?usp=sharing>

## try other model

```

In [11]: # YOLO network hyperparameters
B = 2 # number of bounding box predictions per cell
S = 14 # width/height of network output grid (larger than 7x7 from paper)

learning_rate = 0.001
num_epochs = 20
batch_size = 10 # 24

# Yolo loss component coefficients (as given in Yolo v1 paper)
lambda_coord = 5
lambda_noobj = 0.5

criterion = YoloLoss(S, B, lambda_coord, lambda_noobj)
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate, momentum=

```

```
In [12]: from src.resnet_yolo import resnet101, InceptionV3
         from src.mymodel import FPN, resnet50

         # load_network_path = None #'checkpoints/best_detector.pth'
         # pretrained = True

         # # use to load a previously trained network
         # if load_network_path is not None:
         #     print('Loading saved network from {}'.format(load_network_path))
         #     net = resnet101().to(device)
         #     net.load_state_dict(torch.load(load_network_path))
         # else:
         #     print('Load pre-trained model')
         #     net = resnet101(pretrained=pretrained).to(device)

         # load_network_path = None #'checkpoints/best_detector.pth'
         # pretrained = True

         # use to load a previously trained network
         # if load_network_path is not None:
         #     print('Loading saved network from {}'.format(load_network_path))
         #     net = resnet50().to(device)
         #     net.load_state_dict(torch.load(load_network_path))
         # else:
         #     print('Load pre-trained model')
         #     net = resnet50(pretrained=pretrained).to(device)

         ckpt = torch.load('checkpoints/best_detector.pth')
         net = resnet50()
         net.load_state_dict(ckpt)
         net = net.to(device)

         net = FPN(net)
         net = net.to(device)
```



```

In [13]: best_val_loss = np.inf
learning_rate = 1e-3
for epoch in range(num_epochs):
    torch.cuda.empty_cache()
    net.train()

    # Update learning rate late in training
    if epoch == 30 or epoch == 40:
        learning_rate /= 10.0

    for param_group in optimizer.param_groups:
        param_group['lr'] = learning_rate

    print('\n\nStarting epoch %d / %d' % (epoch + 1, num_epochs))
    print('Learning Rate for this epoch: {}'.format(learning_rate))

    total_loss = collections.defaultdict(int)

    for i, data in enumerate(train_loader):
        data = (item.to(device) for item in data)
        images, target_boxes, target_cls, has_object_map = data
        pred = net(images)
        loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
        for key in loss_dict:
            total_loss[key] += loss_dict[key].item()

        optimizer.zero_grad()
        loss_dict['total_loss'].backward()
        optimizer.step()

        if (i+1) % 50 == 0:
            outstring = 'Epoch [%d/%d], Iter [%d/%d], Loss: ' % ((epoch+1), num_epochs, i+1, 50)
            outstring += ', '.join( "%s=%.3f" % (key[:5], val / (i+1)) for key, val in loss_dict.items())
            print(outstring)

    # evaluate the network on the val data
    if (epoch + 1) % 5 == 0:
        val_aps = evaluate(net, val_dataset_file=annotation_file_val, img_transform=val_transform)
        print(epoch, val_aps)
    with torch.no_grad():
        val_loss = 0.0
        net.eval()
        for i, data in enumerate(val_loader):
            data = (item.to(device) for item in data)
            images, target_boxes, target_cls, has_object_map = data

            pred = net(images)
            loss_dict = criterion(pred, target_boxes, target_cls, has_object_map)
            val_loss += loss_dict['total_loss'].item()
        val_loss /= len(val_loader)

    if best_val_loss > val_loss:
        best_val_loss = val_loss
        print('Updating best val loss: %.5f' % best_val_loss)
        torch.save(net.state_dict(), 'checkpoints/FPN_best_detector.pth')

    if (epoch+1) in [5, 10, 20, 30, 40]:
        torch.save(net.state_dict(), 'checkpoints/FPN_detector_epoch_%d.pth' % (epoch+1))

    torch.save(net.state_dict(), 'checkpoints/FPN_detector.pth')

```

Starting epoch 1 / 20  
Learning Rate for this epoch: 0.001  
Epoch [1/20], Iter [50/376], Loss: total=7.847, reg=3.150, containing\_obj=0.578, no\_obj=0.060, cls=4.059  
Epoch [1/20], Iter [100/376], Loss: total=7.756, reg=3.061, containing\_obj=0.584, no\_obj=0.059, cls=4.052  
Epoch [1/20], Iter [150/376], Loss: total=7.745, reg=3.052, containing\_obj=0.594, no\_obj=0.058, cls=4.042  
Epoch [1/20], Iter [200/376], Loss: total=7.822, reg=3.079, containing\_obj=0.601, no\_obj=0.058, cls=4.084  
Epoch [1/20], Iter [250/376], Loss: total=7.934, reg=3.136, containing\_obj=0.607, no\_obj=0.058, cls=4.133  
Epoch [1/20], Iter [300/376], Loss: total=7.930, reg=3.151, containing\_obj=0.605, no\_obj=0.057, cls=4.117  
Epoch [1/20], Iter [350/376], Loss: total=7.895, reg=3.133, containing\_obj=0.599, no\_obj=0.057, cls=4.106  
Updating best val loss: 8.60175

Starting epoch 2 / 20  
Learning Rate for this epoch: 0.001  
Epoch [2/20], Iter [50/376], Loss: total=7.813, reg=3.137, containing\_obj=0.588, no\_obj=0.056, cls=4.032  
Epoch [2/20], Iter [100/376], Loss: total=8.094, reg=3.237, containing\_obj=0.606, no\_obj=0.055, cls=4.197  
Epoch [2/20], Iter [150/376], Loss: total=7.924, reg=3.174, containing\_obj=0.602, no\_obj=0.056, cls=4.092  
Epoch [2/20], Iter [200/376], Loss: total=7.990, reg=3.194, containing\_obj=0.603, no\_obj=0.056, cls=4.136  
Epoch [2/20], Iter [250/376], Loss: total=7.919, reg=3.158, containing\_obj=0.603, no\_obj=0.056, cls=4.102  
Epoch [2/20], Iter [300/376], Loss: total=8.018, reg=3.193, containing\_obj=0.606, no\_obj=0.056, cls=4.163  
Epoch [2/20], Iter [350/376], Loss: total=8.025, reg=3.204, containing\_obj=0.604, no\_obj=0.056, cls=4.161

Starting epoch 3 / 20  
Learning Rate for this epoch: 0.001  
Epoch [3/20], Iter [50/376], Loss: total=7.754, reg=3.037, containing\_obj=0.603, no\_obj=0.056, cls=4.057  
Epoch [3/20], Iter [100/376], Loss: total=8.046, reg=3.142, containing\_obj=0.623, no\_obj=0.056, cls=4.225  
Epoch [3/20], Iter [150/376], Loss: total=8.239, reg=3.293, containing\_obj=0.616, no\_obj=0.056, cls=4.274  
Epoch [3/20], Iter [200/376], Loss: total=8.253, reg=3.301, containing\_obj=0.618, no\_obj=0.056, cls=4.278  
Epoch [3/20], Iter [250/376], Loss: total=8.143, reg=3.263, containing\_obj=0.612, no\_obj=0.056, cls=4.211  
Epoch [3/20], Iter [300/376], Loss: total=8.128, reg=3.277, containing\_obj=0.608, no\_obj=0.056, cls=4.187  
Epoch [3/20], Iter [350/376], Loss: total=8.027, reg=3.229, containing\_obj=0.605, no\_obj=0.056, cls=4.137  
Updating best val loss: 8.59608

Starting epoch 4 / 20  
Learning Rate for this epoch: 0.001  
Epoch [4/20], Iter [50/376], Loss: total=8.350, reg=3.364, containing\_obj=0.602, no\_obj=0.056, cls=4.327

```
-----
KeyboardInterrupt                                Traceback (most recent call las
/home/vllab/Desktop/A5/A5.ipynb Cell 44 line 2
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=18'>19</a> for i, data in enumerate(train_loader):
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=19'>20</a>     data = (item.to(device) for item in data)
--> <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=20'>21</a>     images, target_boxes, target_cls, has_object
ap = data
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=21'>22</a>     pred = net(images)
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=22'>23</a>     loss_dict = criterion(pred, target_boxes, ta
et_cls, has_object_map)

/home/vllab/Desktop/A5/A5.ipynb Cell 44 line 2
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=16'>17</a> total_loss = collections.defaultdict(int)
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=18'>19</a> for i, data in enumerate(train_loader):
--> <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=19'>20</a>     data = (item.to(device) for item in data)
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=20'>21</a>     images, target_boxes, target_cls, has_object
ap = data
    <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X61sZm
ZQ%3D%3D?line=21'>22</a>     pred = net(images)
```

KeyboardInterrupt:

```
In [14]: net.eval()

# select random image from val set
image_name = random.choice(val_dataset.fnames)
image = cv2.imread(os.path.join(file_root_val, image_name))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

print('predicting...')
result = predict_image(net, image_name, root_img_directory=file_root_val)
for left_up, right_bottom, class_name, _, prob in result:
    color = COLORS[VOC_CLASSES.index(class_name)]
    cv2.rectangle(image, left_up, right_bottom, color, 2)
    label = class_name + str(round(prob, 2))
    text_size, baseline = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX
    p1 = (left_up[0], left_up[1] - text_size[1])
    cv2.rectangle(image, (p1[0] - 2 // 2, p1[1] - 2 - baseline), (p1[0] +
        color, -1)
    cv2.putText(image, label, (p1[0], p1[1] + baseline), cv2.FONT_HERSHEY

plt.figure(figsize = (15,15))
plt.imshow(image)

predicting...
Out[14]: <matplotlib.image.AxesImage at 0x7f19022ad650>
```



```
In [15]: root_test = 'data/VOCdevkit_2007/VOC2007test/JPEGImages/'
file_test = 'data/voc2007test.txt'

ckpt = torch.load('checkpoints/FPN_best_detector.pth')
net = resnet50(pretrained=True).to(device)

net = FPN(net)
net.load_state_dict(ckpt)
net = net.to(device)

preds_submission = test_evaluate(net, test_dataset_file=file_test, img_ro
write_csv(preds_submission)

---Evaluate model on test samples---
28%|██████████          | 1381/4950 [03:06<08:00, 7.42it/s]
```

```

-----
KeyboardInterrupt                                Traceback (most recent call las
/home/vllab/Desktop/A5/A5.ipynb Cell 46 line 1
      <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X63sZ
sZQ%3D%3D?line=7'>8</a> net.load_state_dict(ckpt)
      <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X63sZ
sZQ%3D%3D?line=8'>9</a> net = net.to(device)
--> <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X63sZm
ZQ%3D%3D?line=10'>11</a> preds_submission = test_evaluate(net, test_datas
_file=file_test, img_root=root_test)
      <a href='vscode-notebook-cell:/home/vllab/Desktop/A5/A5.ipynb#X63sZm
ZQ%3D%3D?line=11'>12</a> write_csv(preds_submission)

File ~/Desktop/A5/src/eval_voc.py:207, in test_evaluate(model, test_datas
_file, img_root, test_loader)
    205 model.eval()
    206 for image_path in tqdm(image_list):
--> 207     result = predict_image(model, image_path, root_img_directory=
g_root)
    208     for (
    209         (x1, y1),
    210         (x2, y2),
    (...))
    213         prob,
    214     ) in result: # image_id is actually image_path
    215         preds[class_name].append([image_id, prob, x1, y1, x2, y2])

File ~/Desktop/A5/src/predict.py:144, in predict_image(model, image_name,
oot_img_directory)
    141 img = img.cuda()
    143 pred = model(img) # 1xSxSx(B*5+C)
--> 144 pred = pred.cpu()
    145 boxes, cls_indexs, probs = decoder(pred)
    147 for i, box in enumerate(boxes):

```

KeyboardInterrupt: